


When do true physical inputs help a learned solubility model?

1. The paradox




Premise: truer physical inputs $\rightarrow$ better.

Here, the opposite. So which is wrong — the closure, or the inputs?

2. The measurement

*the true input has an external oracle (VT-2005),
so the true-input error splits exactly:*

$$B = B_{\text{closure}} + B_{\text{insuff}}$$



$$B_{\text{closure}} > B_{\text{insuff}} \quad (P \approx 0.78)$$

**The ceiling is the CLOSURE,
not the inputs.**

3. The map

Two levers flip the sign:
raise closure fidelity, or supervise the latent.

